

# Fundamentals of Database Systems

## COMPSCI 351

Instructor: Sebastian Link

The University of Auckland

# Introduction

## Definition (Database)

A **database** is a very large, integrated and commonly used collection of data

## Definition (Database)

A **database** is a very large, integrated and commonly used collection of data

## Example (Databases model real-world enterprises)

- entities (for example, Rihanna, COMPSCI351)
- relationships (for example, Rihanna is taking COMPSCI351)

# Database (Management) Systems

## Definition (Database Management System)

A **Database Management System** (DBMS) is a

- software package designed to
- store and manage databases

# Database (Management) Systems

## Definition (Database Management System)

A **Database Management System** (DBMS) is a

- software package designed to
- store and manage databases

## Definition (Database System)

A **Database System** (DB) is a DBMS together with a database

- Applications must handle large data sets
  - between main memory and secondary storage
  - for example, buffering, page-oriented access, 32/64-bit addressing, etc.

- Applications must handle large data sets
  - between main memory and secondary storage
  - for example, buffering, page-oriented access, 32/64-bit addressing, etc.
- Special code for different queries



- Applications must handle large data sets
  - between main memory and secondary storage
  - for example, buffering, page-oriented access, 32/64-bit addressing, etc.
- Special code for different queries
- Must protect data from inconsistency due to multiple concurrent users

- Applications must handle large data sets
  - between main memory and secondary storage
  - for example, buffering, page-oriented access, 32/64-bit addressing, etc.
- Special code for different queries
- Must protect data from inconsistency due to multiple concurrent users
- Crash recovery

- Applications must handle large data sets
  - between main memory and secondary storage
  - for example, buffering, page-oriented access, 32/64-bit addressing, etc.
- Special code for different queries
- Must protect data from inconsistency due to multiple concurrent users
- Crash recovery
- Security and access control

# Why Use a DBMS?

- Data independence

# Why Use a DBMS?

- Data independence
- Reduced application development time

# Why Use a DBMS?

- Data independence
- Reduced application development time
- Data integrity and security

# Why Use a DBMS?

- Data independence
- Reduced application development time
- Data integrity and security
- Uniform data administration

# Why Use a DBMS?

- Data independence
- Reduced application development time
- Data integrity and security
- Uniform data administration
- Concurrent access



# Why Use a DBMS?

- Data independence
- Reduced application development time
- Data integrity and security
- Uniform data administration
- Concurrent access
- Recovery from crashes

# Why Use a DBMS?

- Data independence
- Reduced application development time
- Data integrity and security
- Uniform data administration
- Concurrent access
- Recovery from crashes
- Most important factors in designing DBMSs and databases:
  - Efficient access to data
  - Efficient updates of data
  - Usually, there are trade-offs between the two goals

# Why Study Databases?

## Shift from *computation* to *information*

- Always been true for corporate computing
- The Web made this point for personal computing
- More so true for scientific computing

# Why Study Databases?

## Shift from *computation* to *information*

- Always been true for corporate computing
- The Web made this point for personal computing
- More so true for scientific computing

## Big Data (Volume, Variety, Velocity, Veracity)

- Digital libraries, life science projects
- Multimedia databases, social web (Facebook, Twitter, YouTube)
- Semantic annotations, sensor data, uncertain data

# Why Study Databases?

## Shift from *computation* to *information*

- Always been true for corporate computing
- The Web made this point for personal computing
- More so true for scientific computing

## Big Data (Volume, Variety, Velocity, Veracity)

- Digital libraries, life science projects
- Multimedia databases, social web (Facebook, Twitter, YouTube)
- Semantic annotations, sensor data, uncertain data

## DBMS: CompSci as a practical discipline

- AI, languages, logic, multimedia, OS, theory
- Yet traditional focus on real-world applications

# Why Study Databases?

## Shift from *computation* to *information*

- Always been true for corporate computing
- The Web made this point for personal computing
- More so true for scientific computing

## Big Data (Volume, Variety, Velocity, Veracity)

- Digital libraries, life science projects
- Multimedia databases, social web (Facebook, Twitter, YouTube)
- Semantic annotations, sensor data, uncertain data

Need for DBMS  
exploding

## DBMS: CompSci as a practical discipline

- AI, languages, logic, multimedia, OS, theory
- Yet traditional focus on real-world applications

## Definition (Data model)

A collection of concepts for describing data

# Data Models

## Definition (Data model)

A collection of concepts for describing data

## Definition (Schema)

An abstract description of databases using the given data model



# Data Models

## Definition (Data model)

A collection of concepts for describing data

## Definition (Schema)


An abstract description of databases using the given data model

## Example (Relational model of data)

- Most widely used model today
- Main concept: relation, basically a table with rows and columns
- Every relation has a schema, which describes the columns, or fields

# Levels of Abstraction

Many views (external schemata), single logical schema and single physical schema




# Levels of Abstraction

Many views (external schemata), single logical schema and single physical schema


## Views

describe how users see the data



# Levels of Abstraction

Many views (external schemata), single logical schema and single physical schema




Logical schema

defines logical structure

# Levels of Abstraction

Many views (external schemata), single logical schema and single physical schema




## Physical schema

describes the files and  
indexes used

# Levels of Abstraction


Many views (external schemata), single logical schema and single physical schema



Data Definition  
Language (DDL)  
defines database schema

# Levels of Abstraction

Many views (external schemata), single logical schema and single physical schema




Data Manipulation  
Language (DML)

update data in database  
(insert, delete, modify)

# Levels of Abstraction

Many views (external schemata), single logical schema and single physical schema



Query Language  
(QL)

access and retrieve data  
from database



# Levels of Abstraction

Many views (external schemata), single logical schema and single physical schema

## Views


describe how users see the data

## Logical schema

defines logical structure

## Physical schema

describes the files and indexes used



## Data Definition Language (DDL)

defines database schema

## Data Manipulation Language (DML)

update data in database (insert, delete, modify)

## Query Language (QL)

access and retrieve data from database

## Logical schema

- *Students*(sid:string, name:string, login:string, gpa:real)
- *Courses*(cid:string, cname:string, credits:integer)
- *Enrolled*(sid:string, cid:string, grade:string)

# Example: University Database

## External Schema (View)

- *Course\_enrollment*(cid:string,enrollment:integer)
- *Course\_averages*(cid:string,average\_grade:real)

## Logical schema

- *Students*(sid:string, name:string, login:string, gpa:real)
- *Courses*(cid:string, cname:string, credits:integer)
- *Enrolled*(sid:string, cid:string, grade:string)

# Example: University Database

## External Schema (View)

- *Course\_enrollment*(cid:string, enrollment:integer)
- *Course\_averages*(cid:string, average\_grade:real)

## Logical schema

- *Students*(sid:string, name:string, login:string, gpa:real)
- *Courses*(cid:string, cname:string, credits:integer)
- *Enrolled*(sid:string, cid:string, grade:string)

## Physical schema

- Relations stored as unordered files
- Index on first column of *Students*, first of *Courses*, first two of *Enrolled*

# Data Independence: A key benefit of using DBMSs

Applications insulated from how data is structured and stored

# Data Independence: A key benefit of using DBMSs

Applications insulated from how data is structured and stored

## Logical data independence

- external handling of data nearly independent from logical organization
- changes to the logical schema shall not affect the external schema
- only mapping shall be changed
- application shall only see the external schema (impossible in general)

# Data Independence: A key benefit of using DBMSs

Applications insulated from how data is structured and stored

## Logical data independence

- external handling of data nearly independent from logical organization
- changes to the logical schema shall not affect the external schema
- only mapping shall be changed
- application shall only see the external schema (impossible in general)

## Physical data independence

- physical organization of data nearly independent from logical organization
- changes to physical schema have no implications on the logical layer
- can abstract from the realization of the DBMS storage organization
- can reason about data without worrying about physical realization
- can perform physical optimization / tuning

Concurrent execution of user programs is essential for good DBMS performance and business survival



# Concurrency Control

Concurrent execution of user programs is essential for good DBMS performance and business survival

Disk accesses frequent, and relatively slow

important to keep CPU humming by working on several user programs concurrently

# Concurrency Control

Concurrent execution of user programs is essential for good DBMS performance and business survival

Disk accesses frequent, and relatively slow

important to keep CPU humming by working on several user programs concurrently

Interleaving actions of different user programs can lead to inconsistency

e.g., cheque is cleared while account balance is being computed

# Concurrency Control

Concurrent execution of user programs is essential for good DBMS performance and business survival

Disk accesses frequent, and relatively slow

important to keep CPU humming by working on several user programs concurrently

Interleaving actions of different user programs can lead to inconsistency

e.g., cheque is cleared while account balance is being computed

DBMS ensures such problems do not arise

users can pretend they are using a single-user system

# Databases Make These Folks Happy...

End users

get answers to questions they ask

# Databases Make These Folks Happy...

End users

get answers to questions they ask

DBMS vendors

see their products in use

# Databases Make These Folks Happy...

## End users

get answers to questions they ask

## DBMS vendors

see their products in use

## Application programmers

can retrieve and roll out information as required by application

# Databases Make These Folks Happy...

## End users

get answers to questions they ask

## DBMS vendors

see their products in use

## Application programmers

can retrieve and roll out information as required by application

## Database administrators (DBA)

- Designs logical /physical schemas
- Handles security and authorization
- Data availability, crash recovery
- Database tuning as needs evolve

# Databases Make These Folks Happy...

## End users

get answers to questions they ask

## DBMS vendors

see their products in use

## Application programmers

can retrieve and roll out information as required by application

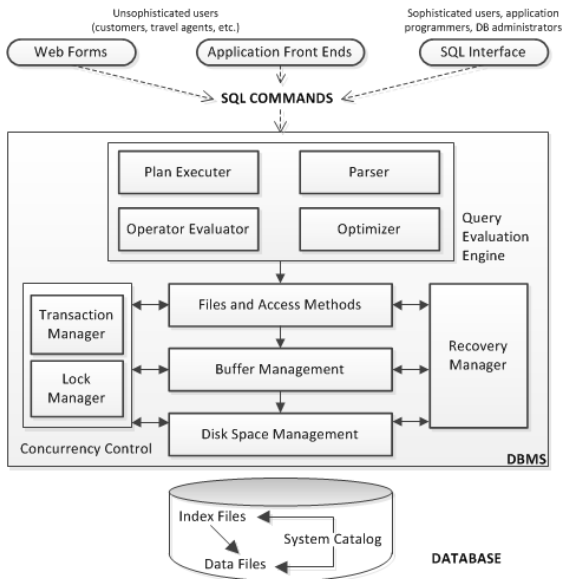
## Database administrators (DBA)

- Designs logical /physical schemas
- Handles security and authorization
- Data availability, crash recovery
- Database tuning as needs evolve

Understanding how a DBMS works is key



# Layered Architecture of a DBMS



# Historical Perspective

- First DBMS was built in the early 1960s
  - based on the network data model

# Historical Perspective


- First DBMS was built in the early 1960s
  - based on the network data model
- IBM developed IMS DBMS in late 1960s
  - based on the hierarchical data model

# Historical Perspective

- First DBMS was built in the early 1960s
  - based on the network data model
- IBM developed IMS DBMS in late 1960s
  - based on the hierarchical data model
- SABRE, an airline reservation system
  - developed by American Airlines and IBM at same time


# Historical Perspective

- First DBMS was built in the early 1960s
  - based on the network data model
- IBM developed IMS DBMS in late 1960s
  - based on the hierarchical data model
- SABRE, an airline reservation system
  - developed by American Airlines and IBM at same time
- Edgar Codd proposed the relational model in 1970s
  - providing a firm foundation for databases
  - won him the Turing Award in 1981




# Historical Perspective

- First DBMS was built in the early 1960s
  - based on the network data model
- IBM developed IMS DBMS in late 1960s
  - based on the hierarchical data model
- SABRE, an airline reservation system
  - developed by American Airlines and IBM at same time
- Edgar Codd proposed the relational model in 1970s
  - providing a firm foundation for databases
  - won him the Turing Award in 1981
- 1980s: SQL developed as part of IBM's system R project



# Historical Perspective

- First DBMS was built in the early 1960s
  - based on the network data model
- IBM developed IMS DBMS in late 1960s
  - based on the hierarchical data model
- SABRE, an airline reservation system
  - developed by American Airlines and IBM at same time
- Edgar Codd proposed the relational model in 1970s
  - providing a firm foundation for databases
  - won him the Turing Award in 1981
- 1980s: SQL developed as part of IBMs system R project
- 1990s: powerful query languages, richer data models (for example, to include images and text)



# Historical Perspective Continued

- Emergence of Enterprise Resource Planning (ERP) and Material Requirements Planning (MRP) packages
  - add a substantial layer on top of usually a relational DBMS
  - e.g., Baan, Oracle, PeopleSoft, SAP, Siebel



# Historical Perspective Continued

- Emergence of Enterprise Resource Planning (ERP) and Material Requirements Planning (MRP) packages
  - add a substantial layer on top of usually a relational DBMS
  - e.g., Baan, Oracle, PeopleSoft, SAP, Siebel
- Web data:
  - originally stored in files,
  - now more common to have an underlying database,
  - DBMS vendors are making DBMSs more web friendly

# Historical Perspective Continued

- Emergence of Enterprise Resource Planning (ERP) and Material Requirements Planning (MRP) packages
  - add a substantial layer on top of usually a relational DBMS
  - e.g., Baan, Oracle, PeopleSoft, SAP, Siebel
- Web data:
  - originally stored in files,
  - now more common to have an underlying database,
  - DBMS vendors are making DBMSs more web friendly
- Today field driven by e-commerce, multimedia databases, digital libraries, life sciences, Web X.0, consolidation of decision making processes and mining of massive data for useful business information

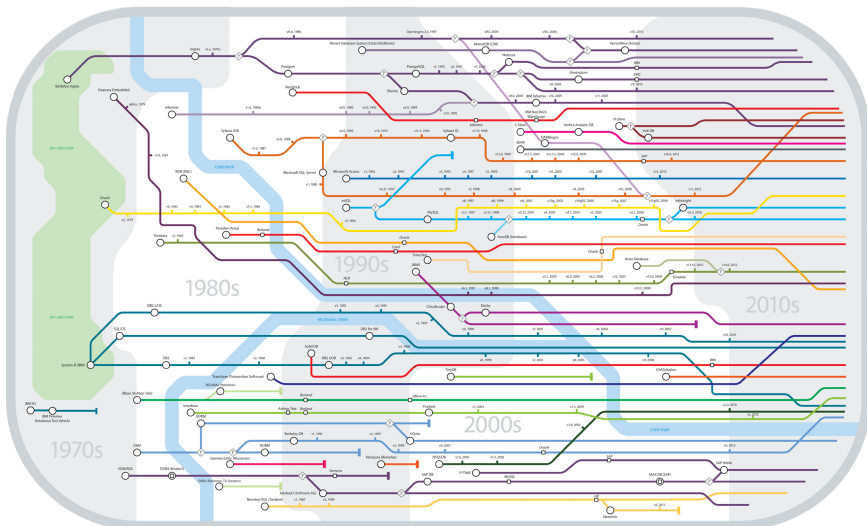
# Historical Perspective Continued

- Emergence of Enterprise Resource Planning (ERP) and Material Requirements Planning (MRP) packages
  - add a substantial layer on top of usually a relational DBMS
  - e.g., Baan, Oracle, PeopleSoft, SAP, Siebel
- Web data:
  - originally stored in files,
  - now more common to have an underlying database,
  - DBMS vendors are making DBMSs more web friendly
- Today field driven by e-commerce, multimedia databases, digital libraries, life sciences, Web X.0, consolidation of decision making processes and mining of massive data for useful business information

Relational databases may not be hot or sexy but for your important data, there is no substitute.

Sean Doherty, Vice-President of EnterpriseDB

# Genealogy of Relational Database Management Systems



# Summary

- Value of data to an organization widely recognized
- Amount and type of data used is exploding
- Data can become a liability rather than an asset if not managed
- DBMS used to maintain and query large data sets
- Benefits include data integrity and security, quick application development, concurrent access and recovery from system crashes
- Overriding factors in database design are efficient access to and updates of data
- Levels of abstraction give data independence
- A DBMS typically has a layered architecture
- DBAs hold responsible jobs and are well-paid
- DBMS R&D broad and exciting with deep interplay between theory and practice